

HEART STROKE PREDICTION SYSTEM

In this machine learning project, the overall topic that will be resolved is in the field of stroke health, where it will try to predict the possibility of a stroke in a person with certain conditions based on several factors including: age, certain diseases (hypertension, heart disease), smoking, etc.

Install and import required libraries.

```
# library for data processing
import pandas as pd
from sklearn.preprocessing import LabelEncoder
#Library for data visualization
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
# library for modeling
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
# library for model evaluation
from sklearn.metrics import accuracy_score
```

Read data with pandas

```
df = pd.read_csv('healthcare-dataset-stroke-data.csv')
```

Explore Dataset information

```
df.tail()
```

	id	gender	age	hypertension	heart_disease	ever_married	\
5105	18234	Female	80.0	1	0	Yes	
5106	44873	Female	81.0	0	0	Yes	
5107	19723	Female	35.0	0	0	Yes	
5108	37544	Male	51.0	0	0	Yes	
5109	44679	Female	44.0	0	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	
smoking_status	\				
5105	Private	Urban	83.75	NaN	never

smoked					
5106	Self-employed	Urban	125.20	40.0	never
smoked					
5107	Self-employed	Rural	82.99	30.6	never
smoked					
5108	Private	Rural	166.29	25.6	formerly
smoked					
5109	Govt_job	Urban	85.28	26.2	
Unknown					

	stroke
5105	0
5106	0
5107	0
5108	0
5109	0

check dataset info

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   id                    5110 non-null   int64  
 1   gender                5110 non-null   object  
 2   age                  5110 non-null   float64 
 3   hypertension          5110 non-null   int64  
 4   heart_disease         5110 non-null   int64  
 5   ever_married          5110 non-null   object  
 6   work_type             5110 non-null   object  
 7   Residence_type        5110 non-null   object  
 8   avg_glucose_level     5110 non-null   float64 
 9   bmi                   4909 non-null   float64 
10   smoking_status        5110 non-null   object  
11   stroke                5110 non-null   int64  
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
```

describe numeric column

```
df.describe()
```

	id	age	hypertension	heart_disease	\
count	5110.000000	5110.000000	5110.000000	5110.000000	
mean	36517.829354	43.226614	0.097456	0.054012	
std	21161.721625	22.612647	0.296607	0.226063	
min	67.000000	0.080000	0.000000	0.000000	
25%	17741.250000	25.000000	0.000000	0.000000	
50%	36932.000000	45.000000	0.000000	0.000000	
75%	54682.000000	61.000000	0.000000	0.000000	
max	72940.000000	82.000000	1.000000	1.000000	

	avg_glucose_level	bmi	stroke
count	5110.000000	4909.000000	5110.000000
mean	106.147677	28.893237	0.048728
std	45.283560	7.854067	0.215320
min	55.120000	10.300000	0.000000
25%	77.245000	23.500000	0.000000
50%	91.885000	28.100000	0.000000
75%	114.090000	33.100000	0.000000
max	271.740000	97.600000	1.000000

Check value counts

```
df["stroke"].value_counts()
```

```
0    4861
```

```
1     249
```

```
Name: stroke, dtype: int64
```

```
def get_smoke_count(smoking_status):
    mask=df['smoking_status']==smoking_status
    return df[mask]
```

```
get_smoke_count('Unknown')
```

	id	gender	age	hypertension	heart_disease	ever_married
work_type \						
8 Private	27419	Female	59.0	0	0	Yes
9 Private	60491	Female	78.0	0	0	Yes
13 Private	8213	Male	78.0	0	1	Yes
19 Govt_job	25226	Male	57.0	0	1	No
23 Private	64778	Male	82.0	0	1	Yes
...
5098	579	Male	9.0	0	0	No

```

children
5101 36901 Female 45.0 0 0 Yes
Private
5103 22127 Female 18.0 0 0 No
Private
5104 14180 Female 13.0 0 0 No
children
5109 44679 Female 44.0 0 0 Yes
Govt_job

```

```

      Residence_type avg_glucose_level  bmi smoking_status  stroke
8              Rural          76.15   NaN      Unknown      1
9              Urban          58.57  24.2      Unknown      1
13             Urban         219.84   NaN      Unknown      1
19             Urban         217.08   NaN      Unknown      1
23             Rural         208.30  32.5      Unknown      1
...             ...             ...   ...             ...
5098            Urban          71.88  17.5      Unknown      0
5101            Urban          97.95  24.5      Unknown      0
5103            Urban          82.85  46.9      Unknown      0
5104            Rural         103.08  18.6      Unknown      0
5109            Urban          85.28  26.2      Unknown      0

```

```
[1544 rows x 12 columns]
```

```
df["smoking_status"].value_counts()
```

```

never smoked    1892
Unknown         1544
formerly smoked    885
smokes           789
Name: smoking_status, dtype: int64

```

Check for null values

```
df.isnull().sum()
```

```

id              0
gender          0
age            0
hypertension    0
heart_disease   0
ever_married    0
work_type       0
Residence_type  0
avg_glucose_level 0
bmi            201
smoking_status  0

```

```

stroke          0
dtype: int64

null_values = df['bmi'].isnull().sum()
null_values

201

null_percentage = df['bmi'].isnull().mean()*100
null_percentage

3.9334637964774952

column_mean = df['bmi'].mean()
df['bmi'].fillna(column_mean, inplace=True)

df.columns

Index(['id', 'gender', 'age', 'hypertension', 'heart_disease',
      'ever_married',
      'work_type', 'Residence_type', 'avg_glucose_level', 'bmi',
      'smoking_status', 'stroke'],
      dtype='object')

columns = ['gender', 'ever_married', 'work_type',
          'Residence_type', 'smoking_status']

for col in columns:
    unique_values = df[col].unique()
    print(f"Unique values for{col}:{unique_values}")

Unique values forgender:['Male' 'Female' 'Other']
Unique values forever_married:['Yes' 'No']
Unique values forwork_type:['Private' 'Self-employed' 'Govt_job'
                             'children' 'Never_worked']
Unique values forResidence_type:['Urban' 'Rural']
Unique values forsmoking_status:['formerly smoked' 'never smoked'
                                  'smokes' 'Unknown']

df['gender'].value_counts()

Female    2994
Male      2115
Other         1
Name: gender, dtype: int64

# In gender column there is only one value for other so we replace it
with male
df['gender'] = df['gender'].replace('Other', 'Male')

df['gender'].value_counts()

```

```

Female      2994
Male        2116
Name: gender, dtype: int64

mask1 = df[['gender', 'ever_married', 'work_type',
            'Residence_type', 'smoking_status']].nunique()
mask1

gender      2
ever_married 2
work_type   5
Residence_type 2
smoking_status 4
dtype: int64

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    5110 non-null   int64
1   gender                5110 non-null   object
2   age                   5110 non-null   float64
3   hypertension          5110 non-null   int64
4   heart_disease         5110 non-null   int64
5   ever_married          5110 non-null   object
6   work_type             5110 non-null   object
7   Residence_type        5110 non-null   object
8   avg_glucose_level     5110 non-null   float64
9   bmi                   5110 non-null   float64
10  smoking_status        5110 non-null   object
11  stroke                 5110 non-null   int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB

```

Data Visualisation


```

corr_matrix = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, vmin=-1, vmax=1, cmap='coolwarm')
plt.show()

```

C:\Users\prita\AppData\Local\Temp\ipykernel_17232\1618691362.py:1:
FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corr_matrix = df.corr()
C:\Users\prita\AppData\Local\Temp\ipykernel_17232\1618691362.py:3:
FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only
valid columns or specify the value of numeric_only to silence this
warning.
sns.heatmap(df.corr(), annot=True, vmin=-1, vmax=1, cmap='coolwarm')
```



Feature Engineering

```
# Creating age group categories

print(f'maximum age variable: {df["age"].max()}')
print(f'minimum age variable: {df["age"].min()}')
print(f'Number of age variable: {df["age"].nunique()}')
```

```
maximum age variable: 82.0
minimum age variable: 0.08
Number of age variable: 104
```

```
# collapse age group categories
```

```
ranges = [0,13,18,45,60,100]
group_names = ['Children', 'Teens', 'Adults', 'Mid-adults', 'Elderly']
df['age_group'] = pd.cut(df['age'],bins=ranges,labels=group_names)
df['age_group'].unique()
```

```
['Elderly', 'Mid-adults', 'Adults', 'Children', 'Teens']
Categories (5, object): ['Children' < 'Teens' < 'Adults' < 'Mid-
adults' < 'Elderly']
```

```
# For BMI
```

```
print(f'maximum age variable: {df["bmi"].max()}')
print(f'minimum age variable: {df["bmi"].min()}')
print(f'Number of age variable: {df["bmi"].nunique()}')
```

```
maximum age variable: 97.6
minimum age variable: 10.3
Number of age variable: 419
```

```
# collapse bmi into fewer groups
```

```
ranges = [0,19,25,30,100]
group_names = ['Underweight', 'Normal', 'Overweight', 'Obesity']
df['bmi_group'] = pd.cut(df['bmi'],bins=ranges,labels=group_names)
df['bmi_group'].unique()
```

```
['Obesity', 'Overweight', 'Normal', 'Underweight']
Categories (4, object): ['Underweight' < 'Normal' < 'Overweight' <
'Obesity']
```

```
# for avg glucose level
```

```
print(f'maximum age variable: {df["avg_glucose_level"].max()}')
print(f'minimum age variable: {df["avg_glucose_level"].min()}')
print(f'Number of age variable: {df["avg_glucose_level"].nunique()}')
```

```
maximum age variable: 271.74
minimum age variable: 55.12
Number of age variable: 3979
```


```
ranges = [0, 70, 99, 125, 280]
group_names = ['Low', 'Normal', 'High', 'Very_high']
df['avg_glucose_level_group'] =
pd.cut(df['avg_glucose_level'],bins=ranges,labels=group_names)
df['avg_glucose_level_group'].unique()
```



```
['Very_high', 'High', 'Normal', 'Low']  
Categories (4, object): ['Low' < 'Normal' < 'High' < 'Very_high']
```

Exploratory data analysis(EDA)

```
sns.countplot(x='stroke',data=df)  
plt.show()
```



```
df.head()
```


	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	\
0	Private	Urban	228.69	36.600000	
1	Self-employed	Rural	202.21	28.893237	
2	Private	Rural	105.92	32.500000	
3	Private	Urban	171.23	34.400000	

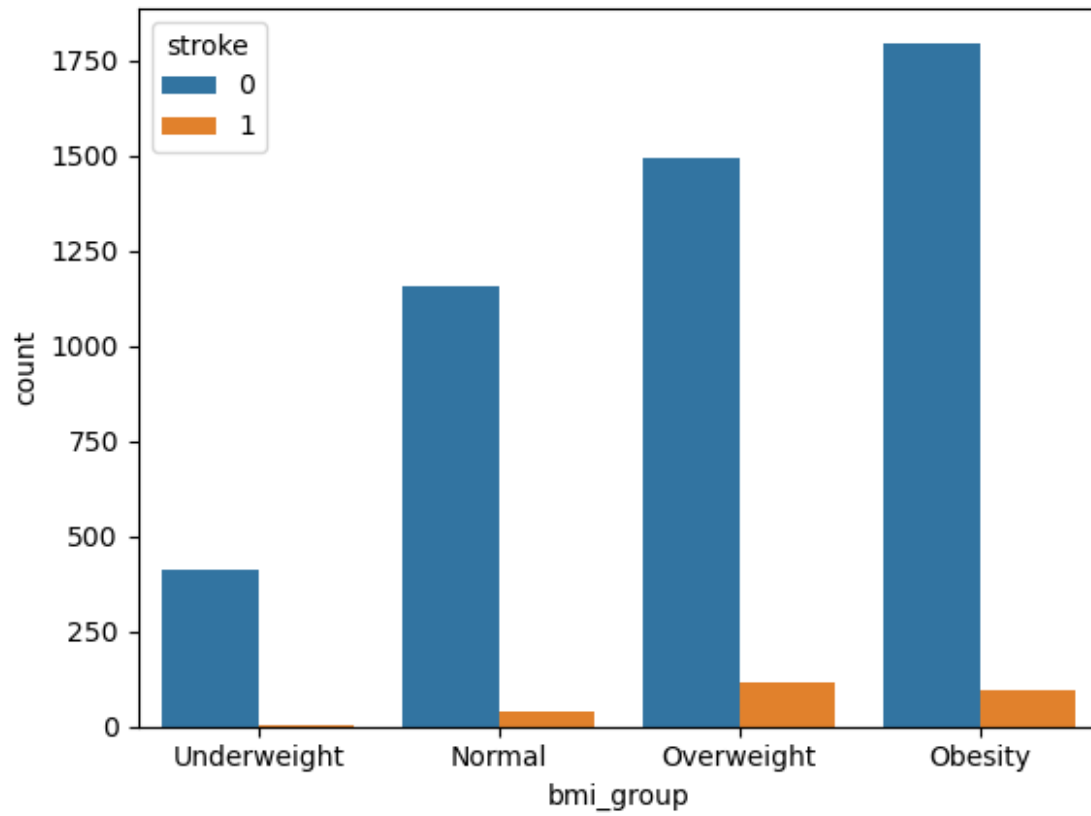
```
4 Self-employed Rural 174.12 24.000000
```

	smoking_status	stroke	age_group	bmi_group
0	formerly smoked	1	Elderly	Obesity
1	never smoked	1	Elderly	Overweight
2	never smoked	1	Elderly	Obesity
3	smokes	1	Mid-adults	Obesity
4	never smoked	1	Elderly	Normal


```
sns.countplot(x='age_group', hue='stroke', data=df)
plt.show()
```




```
sns.countplot(x='bmi_group', hue='stroke', data=df)
plt.show()
```



```
sns.countplot(x='avg_glucose_level_group',hue='stroke',data=df)  
plt.show()
```



```
sns.barplot(x='heart_disease',y='age',data=df,hue='stroke')  
<Axes: xlabel='heart_disease', ylabel='age'>
```



df

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	
...
5105	18234	Female	80.0	1	0	Yes	
5106	44873	Female	81.0	0	0	Yes	
5107	19723	Female	35.0	0	0	Yes	
5108	37544	Male	51.0	0	0	Yes	
5109	44679	Female	44.0	0	0	Yes	


	work_type	Residence_type	avg_glucose_level	bmi	\
0	Private	Urban	228.69	36.600000	
1	Self-employed	Rural	202.21	28.893237	
2	Private	Rural	105.92	32.500000	
3	Private	Urban	171.23	34.400000	
4	Self-employed	Rural	174.12	24.000000	
...
5105	Private	Urban	83.75	28.893237	
5106	Self-employed	Urban	125.20	40.000000	

5107	Self-employed	Rural	82.99	30.600000
5108	Private	Rural	166.29	25.600000
5109	Govt_job	Urban	85.28	26.200000


	smoking_status	stroke	age_group	bmi_group
avg_glucose_level_group				
0	formerly smoked	1	Elderly	Obesity
Very_high				
1	never smoked	1	Elderly	Overweight
Very_high				
2	never smoked	1	Elderly	Obesity
High				
3	smokes	1	Mid-adults	Obesity
Very_high				
4	never smoked	1	Elderly	Normal
Very_high				
...
...				
5105	never smoked	0	Elderly	Overweight
Normal				
5106	never smoked	0	Elderly	Obesity
Very_high				
5107	never smoked	0	Adults	Obesity
Normal				
5108	formerly smoked	0	Mid-adults	Overweight
Very_high				
5109	Unknown	0	Adults	Overweight
Normal				

[5110 rows x 15 columns]


```
sns.countplot(x='gender',hue='stroke',data=df)
plt.show()
```



```
stroke = df[df['stroke']==1]
sns.displot(stroke['avg_glucose_level'], kind='kde')
sns.set_style('darkgrid')
plt.show()
```



```
stroke = df[df['stroke']==1]
sns.displot(stroke['bmi'], kind='kde')
sns.set_style('darkgrid')
plt.show()
```

Preprocessing

Binary Encoding

```
#Instantiate LabelEncoder
labelencoder = LabelEncoder()

#Binary Encoding(encoding object columns with 2 unique values)
binary_cols = ['ever_married', 'Residence_type', 'gender']
for col in binary_cols:
    df[col]=labelencoder.fit_transform(df[col])
```

Label Encoding

```
categorical_cols = ['age_group', 'bmi_group',
                    'avg_glucose_level_group']

label_encoder = LabelEncoder()
for col in categorical_cols:
    df[col] = label_encoder.fit_transform(df[col])
```

One hot encoding

#Encode object columns that more than 2 unique values

```
df = pd.get_dummies(df, columns=['work_type', 'smoking_status'],  
drop_first=True)
```

df

	id	gender	age	hypertension	heart_disease	
ever_married	\					
0	9046	1	67.0	0	1	1
1	51676	0	61.0	0	0	1
2	31112	1	80.0	0	1	1
3	60182	0	49.0	0	0	1
4	1665	0	79.0	1	0	1
...
5105	18234	0	80.0	1	0	1
5106	44873	0	81.0	0	0	1
5107	19723	0	35.0	0	0	1
5108	37544	1	51.0	0	0	1
5109	44679	0	44.0	0	0	1
	Residence_type	avg_glucose_level	bmi	stroke	age_group	
	\					
0	1	228.69	36.600000	1	2	
1	0	202.21	28.893237	1	2	
2	0	105.92	32.500000	1	2	
3	1	171.23	34.400000	1	3	
4	0	174.12	24.000000	1	2	
...
5105	1	83.75	28.893237	0	2	
5106	1	125.20	40.000000	0	2	
5107	0	82.99	30.600000	0	0	

5108	0	166.29	25.600000	0	3
------	---	--------	-----------	---	---

5109	1	85.28	26.200000	0	0
------	---	-------	-----------	---	---

	bmi_group	avg_glucose_level_group	work_type_Never_worked	\
0	1	3	0	
1	2	3	0	
2	1	0	0	
3	1	3	0	
4	0	3	0	
...	
5105	2	2	0	
5106	1	3	0	
5107	1	2	0	
5108	2	3	0	
5109	2	2	0	

	work_type_Private	work_type_Self-employed	work_type_children	\
0	1	0	0	
1	0	1	0	
2	1	0	0	
3	1	0	0	
4	0	1	0	
...	
5105	1	0	0	
5106	0	1	0	
5107	0	1	0	
5108	1	0	0	
5109	0	0	0	

	smoking_status_formerly smoked	smoking_status_never smoked	\
0	1	0	
1	0	1	
2	0	1	
3	0	0	
4	0	1	
...	

5105	0	1
5106	0	1
5107	0	1
5108	1	0
5109	0	0

	smoking_status_smokes
0	0
1	0
2	0
3	1
4	0
...	...
5105	0
5106	0
5107	0
5108	0
5109	0

[5110 rows x 20 columns]

Pre Modeling Steps

```
# separate feature and target
X = df.drop('stroke', axis=1)
y = df['stroke']

# using SMOTE Techniqe
# sm = SMOTE(random_state=111)
# X_sm , y_sm = sm.fit_resample(X,y)
#print(f'''Shape of X before SMOTE:{X.shape}
#Shape of X after SMOTE:{X_sm.shape}''', "\n\n")

# print(f'''Target Class distributuion before SMOTE:\
n{y.value_counts(normalize=True)}
# Target Class distributuion after SMOTE :\
n{y_sm.value_counts(normalize=True)}''')
```

Split train data and test data

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Machine Learning Modeling

1) Random forest Classifier

```
# Create simple model
rf_classifier = RandomForestClassifier(n_estimators=100,
random_state=42)

rf_classifier.fit(X_train, y_train)

RandomForestClassifier(random_state=42)

# Test model with test data
y_pred = rf_classifier.predict(X_test)

# Simple model report
rf_accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {rf_accuracy}")

Accuracy: 0.9393346379647749
```

2) Logistic Regression

```
logistic_reg = LogisticRegression()
<IPython.core.display.Javascript object>
logistic_reg.fit(X_train, y_train)
LogisticRegression()

# Test model with test data
y_pred = logistic_reg.predict(X_test)

# Simple model report
log_reg_accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {log_reg_accuracy}")

Accuracy: 0.9373776908023483
```

3) Decision Tree

```
# Create the Decision Tree Classifier
decision_tree_classifier = DecisionTreeClassifier()

# Train the classifier on the training data
decision_tree_classifier.fit(X_train, y_train)

DecisionTreeClassifier()
```

```
# Make predictions on the test set  
y_pred = decision_tree_classifier.predict(X_test)
```

```
# Calculate the accuracy of the classifier  
dt_accuracy = accuracy_score(y_test, y_pred)  
print(f"Accuracy: {dt_accuracy}")
```

```
Accuracy: 0.9158512720156555
```